Felipe P. dos Santos

Senior AI Engineer | LLM Platforms, RAG, Agents & Applied GenAI
Curitiba, Brazil | santospefelipe@gmail.com | mefelipe.com.br | LinkedIn: felipepesantos

---

Summary

Senior AI Engineer who takes applied GenAI to production. I own the shared LLM platform behind the AI products of Unimed Goiânia, a Brazilian health insurer with ~400k members: open-source models on vLLM over AWS GPUs, an OpenAI-compatible gateway, observability, and model fallback. Earlier I built RAG and multi-tool agents at Grupo Boticário, led engineering at a startup, and founded a B2B SaaS (7,000+ users).

Experience

Unimed Goiânia (via DeepDiveIA) | Senior AI Engineer (Nov 2026 – Present)

- **LLM Platform:** Own the shared LLM platform behind the AI products: an OpenAI-compatible FastAPI gateway (per-app keys, quotas, model fallback) over a multimodal open-source model on vLLM (AWS GPU), hybrid search on pgvector, and per-call cost tracking in Langfuse and Grafana.
- **Architecture & Technical Review:** Authored the target architecture (local-first, in-country paid fallback) and the platform's technical review: model selection for one 24 GB GPU, security against the OWASP Top 10 for LLMs, LGPD for health data, and a prioritized pre-production roadmap.
- **Product Integration:** Support the teams integrating the platform into three healthcare AI products: a chat assistant that analyzes Tableau dashboards, legal-claim value estimation from poor-quality scans (vision/OCR), and medical-request transcription into procedure codes.

Grupo Boticário | AI / GenAI & Software Engineer (Mid-level) (Jan 2026 – Nov 2026)

- **RAG & Multi-tool Agents:** Designed a RAG pipeline over support materials, trainings, and operational guidance (vector database, grounded retrieval) and implemented agents with multiple tools (database lookup, ticket/board actions, Slack) inside a documented harness.
- **AI Multiplier & Enablement:** Selected to multiply AI capability across the company: trained other teams, coached GenAI implementation best practices, and unblocked squads during adoption.
- **App Super (10,000+ employees):** Engineer on a workforce-management PWA (React, Node.js, MongoDB, SSO, AWS EKS) that replaced dozens of spreadsheets: +20% productivity for plant coordinators and −10% team-allocation time, on an operational cost base of roughly R$63M/year.

ASSUMTEK (startup) | Full-Stack Engineer (Mid-level) & Tech Lead (Jan 2025 – Jan 2026)

- **Tech Lead & HR Platform:** Ran demand management, solution architecture, and 1:1s. Shipped a greenfield system (Node, Postgres, Next.js, Vercel, Railway) for contractor time tracking and payroll with complex rules (hour caps, approvals, dynamic pay), cutting HR payroll close time by 70%.
- **Automation & BI:** Built integrations (Node.js, n8n, Make) for Meta Ads, CRMs, and internal data, and Node.js/Python pipelines feeding Power BI for Product, Marketing, IT, HR, and B2B clients.

Sentier | Founder & Full-Stack Engineer (B2B & SaaS) (Jan 2021 – Jan 2025)

- **Multi-tenant Barbershop SaaS:** Built scheduling, staff payroll, subscriptions, and automated reminders (Next.js, Node.js, Prisma, AWS S3, Resend, Vercel). 7,000+ registered users; clients saved about 17–20% of daily operational time.
- **End-to-End Product Delivery:** Technical lead and primary engineer on 40+ complete B2B projects (custom systems, landing pages, automations, and BI), from Figma to production.

Skills

AI & LLMs: RAG, embeddings / vector stores (pgvector), LangChain, multi-tool agents, prompt & context engineering, OpenAI SDK, Gemini SDK, AWS Bedrock

LLM Serving & Ops: vLLM, open-source LLMs (Qwen, Gemma), OpenAI-compatible gateways, model fallback, observability (Langfuse, Grafana), LLM security (OWASP), LGPD

Software: Python, FastAPI, Node.js, TypeScript/JavaScript, React, Next.js, Express, Prisma, PostgreSQL, MongoDB, REST APIs, DDD

Cloud & Delivery: AWS (S3, SQS, Lambda, EKS), Docker, CI/CD, Azure DevOps, Vercel, Railway, New Relic, Git Flow

Automation & Other: Make, n8n, Power BI, Figma, Scrum

In progress: AWS Bedrock, Anthropic AI certifications

Education & courses

Academic Background

Analysis and Systems Development | FAEL (Jan 2021 – Dec 2025)

Specializations and Courses

- Desenvolve Program (Full Stack, Node, React, Next, AI) | Grupo Boticário (2025)
- Business Fundamentals | Harvard Business Education (2025)
- Data Analysis & BI | Google (2025)
- Computer Science, CS50x | Harvard University (2021)
- AWS Academy Cloud Foundations | AWS (2021)

Volunteering & social impact

Active Global Shapers (World Economic Forum) volunteer training public-school teachers on practical AI and child-protection awareness; regional leader and IT volunteer for The Church of Jesus Christ of Latter-day Saints (2023–2024); community work with oncology patients plus web/AV support for NGOs.